

- When conducting the work of Lab 11, we conducted the test that uses the Central Limit Theorem even though the sample size was “small” (i.e., $n < 30$). It turns out, that how “far off” the t -test is can be computed using a first-order Edgeworth approximation for the error. Below, we will do this for the the further observations.

(a) Boos and Hughes-Oliver (2000) note that

$$P(T \leq t) \approx F_Z(t) + \underbrace{\frac{\text{skew}}{\sqrt{n}} \frac{(2t^2 + 1)}{6} f_Z(t)}_{\text{error}},$$

where $f_Z(\cdot)$ and $F_Z(\cdot)$ are the Gaussian PDF and CDF and skew is the skewness of the data. What is the potential error in the computation of the p -value when testing $H_0 : \mu_X = 0; H_a : \mu_X < 0$ using the zebra finch further data?

```
library(tidyverse)
library(e1071)
dat1 = read_csv("zebrafinches.csv")

## Rows: 25 Columns: 3
## -- Column specification -----
## Delimiter: ","
## dbl (3): closer, further, diff
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.

dat = dat1$`further`
t.val = as.numeric(t.test(dat, alternative = "less", mu = 0)$statistic)
cdf.val = pnorm(t.val)
pdf.val = dnorm(t.val)
skew = skewness(dat)
n = length(dat)
(error = cdf.val + (skew/sqrt(n))*(2*(t.val)^2+1)/6*pdf.val))

## [1] -1.189164e-13
```

- (b) Compute the error for t statistics from -10 to 10 and plot a line that shows the error across t . Continue to use the skewness and the sample size for the zebra finch further data.

```
total.error = vector()
seq = seq(-10,10,.01)
for(t in seq){
  curr.cdf.val = pnorm(t)
  curr.pdf.val = dnorm(t)
  skew = skewness(dat)
  n = length(dat)
  curr.error = (skew/sqrt(n))*(2*(t)^2+1)/6*curr.pdf.val
  total.error = append(total.error, curr.error)
}
total.error = tibble(x = total.error) |>
mutate(seq)
```

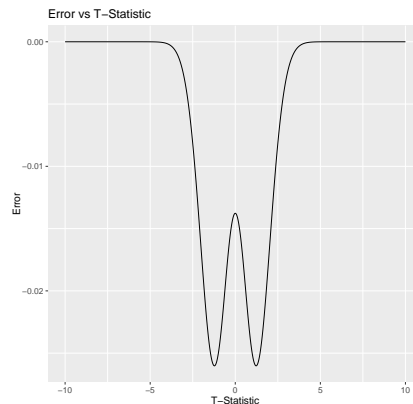


Figure 1:

- (c) Suppose we wanted to have a tail probability within 10% of the desired $\alpha = 0.05$. Recall we did a left-tailed test using the further data. How large of a sample size would we need? That is, we need to solve the error formula equal to 10% of the desired left-tail probability:

$$0.10\alpha \stackrel{set}{=} \underbrace{\frac{\text{skew}}{\sqrt{n}} \frac{(2t^2 + 1)}{6} f_Z(t)}_{\text{error}},$$

which yields

$$n = \left(\frac{\text{skew}}{6(0.10\alpha)} (2t^2 + 1) f_Z(t) \right)^2.$$

```
t.val.c = qnorm(.05)
pdf.val.c = dnorm(t.val.c)
(n.c = (skew/(6*(.1*.05)))*(2*t.val.c^2+1)*pdf.val.c)^2

## [1] 520.8876
```

2. Complete the following steps to revisit the analyses from lab 11 using the bootstrap procedure.

- (a) Now, consider the zebra finch data. We do not know the generating distributions for the closer, further, and difference data, so perform resampling to approximate the sampling distribution of the T statistic:

$$T = \frac{\bar{x}_r - 0}{s/\sqrt{n}},$$

where \bar{x}_r is the mean computed on the r^{th} resample and s is the sample standard deviation from the original samples. At the end, create an object called `resamples.null.closer`, for example, and store the resamples shifted to ensure they are consistent with the null hypotheses at the average (i.e., here ensure the shifted resamples are 0 on average, corresponding to $t = 0$, for each case).

```
#Farther
t.dist.dat = vector()
s = sd(dat)
for(i in 1:10000){
  curr.val = (mean(sample(dat, replace = TRUE)) - 0)/(s/sqrt(n))
  t.dist.dat = append(t.dist.dat, curr.val)
}
ncp = -1*(mean(t.dist.dat))
resamples.null.farther = t.dist.dat + ncp
(mean(resamples.null.farther))

## [1] 3.934297e-16

#Closer
t.close.dat = vector()
s = sd(dat1$closer)
for(i in 1:10000){
  curr.val = (mean(sample(dat1$closer, replace = TRUE)) - 0)/(s/sqrt(n))
  t.close.dat = append(t.close.dat, curr.val)
}
ncp2 = -1*(mean(t.close.dat))
resamples.null.closer = t.close.dat + ncp2
(mean(resamples.null.closer))

## [1] 5.778711e-16

#Difference
t.diff.dat = vector()
s = sd(dat1$diff)
for(i in 1:10000){
  curr.val = (mean(sample(dat1$diff, replace = TRUE)) - 0)/(s/sqrt(n))
  t.diff.dat = append(t.diff.dat, curr.val)
}
ncp3 = -1*(mean(t.diff.dat))
resamples.null.diff = t.diff.dat + ncp3
(mean(resamples.null.diff))

## [1] 8.424261e-16
```

- (b) Compute the bootstrap p -value for each test using the shifted resamples. How do these compare to the t -test p -values?

```
#Farther p-value
x.bar.far = mean(t.dist.dat)
(p.val.far = mean(resamples.null.farther <= x.bar.far))

## [1] 0

#Closer p-value
x.bar.close = mean(t.close.dat)
(p.val.close = mean(resamples.null.closer >= x.bar.close))

## [1] 0

#Difference p-value
x.bar.diff = mean(t.diff.dat)
x.bar.diff.lower = 0-x.bar.diff
x.bar.diff.upper = 0+x.bar.diff

(p.val.diff = mean(resamples.null.diff <= x.bar.diff.lower) + mean(resamples.null.diff >= x.bar.diff.upper))

## [1] 0
```

- (c) What is the 5th percentile of the shifted resamples under the null hypothesis? Note this value approximates $t_{0.05, n-1}$. Compare these values in each case.

```
#Farther
quantile(resamples.null.farther, prob = .025)

##      2.5%
## -2.02765

#Closer
quantile(resamples.null.closer, prob = .025)

##      2.5%
## -1.902959

#Difference
quantile(resamples.null.diff, prob = .025)

##      2.5%
## -1.834268
```

- (d) Compute the bootstrap confidence intervals using the resamples. How do these compare to the t -test confidence intervals?

```
#Farther
n = length(dat1$farther)
sd.f = sd(dat1$farther)
lower.f = quantile(t.dist.dat, prob = .025)
x.bar.lower.f = lower.f/sqrt(n)*sd.f
upper.f = quantile(t.dist.dat, prob = .975)
x.bar.upper.f = upper.f/sqrt(n)*sd.f

c(x.bar.lower.f, x.bar.upper.f)

##      2.5%      97.5%
## -0.2556874 -0.1554911

#Closer
sd.c = sd(dat1$closer)
lower.c = quantile(t.close.dat, prob = .025)
x.bar.lower.c = lower.c/sqrt(n)*sd.c
upper.c = quantile(t.close.dat, prob = .975)
x.bar.upper.c = upper.c/sqrt(n)*sd.c

c(x.bar.lower.c, x.bar.upper.c)

##      2.5%      97.5%
## 0.1200736 0.1926468

#Difference
sd.d = sd(dat1$diff)
lower.d = quantile(t.diff.dat, prob = .025)
x.bar.lower.d = lower.d/sqrt(n)*sd.d
upper.d = quantile(t.diff.dat, prob = .975)
x.bar.upper.d = upper.d/sqrt(n)*sd.d

c(x.bar.lower.d, x.bar.upper.d)
```

```
##      2.5%      97.5%
## 0.2810613 0.4434062
```

3. Complete the following steps to revisit the analyses from lab 11 using the randomization procedure.

- (a) Now, consider the zebra finch data. We do not know the generating distributions for the closer, further, and difference data, so perform the randomization procedure

```
#####
# Further Data
#####
mu0 = 0
R = 10000
rand = tibble(xbars = rep(NA, 10000))
samp.mean = mean(dat1$further)
x.shift = dat1$further - mu0
#Doing the randomization process
for(i in 1:R){
  curr.samp = x.shift * sample(x = c(-1,1),
                              size = length(x.shift),
                              replace = T)

  rand$xbars[i] = mean(curr.samp)
}
#Shifting back
rand1 = rand |>
  mutate(xbars = xbars + mu0)
#####
# Closer Data
#####
rand.cl = tibble(xbars = rep(NA, 10000))
samp.mean.cl = mean(dat1$closer)
x.shift = dat1$closer - mu0
#Doing the randomization process
for(i in 1:R){
  curr.samp = x.shift * sample(x = c(-1,1),
                              size = length(x.shift),
                              replace = T)

  rand.cl$xbars[i] = mean(curr.samp)
}
#Shifting back
rand.cl = rand.cl |>
  mutate(xbars = xbars + mu0)
#####
# Different Data
#####

rand.diff = tibble(xbars = rep(NA, 10000))
samp.mean.diff = mean(dat1$diff)
x.shift = dat1$diff - mu0
#Doing the randomization process
for(i in 1:R){
  curr.samp = x.shift * sample(x = c(-1,1),
                              size = length(x.shift),
                              replace = T)

  rand.diff$xbars[i] = mean(curr.samp)
}
#Shifting back
rand.diff = rand.diff |>
  mutate(xbars = xbars + mu0)
```

- (b) Compute the randomization test p -value for each test.

```
#Farther
(furth.p = mean(rand1<=samp.mean))

## [1] 0

#Closer
(closer.p = mean(rand>=samp.mean.cl))

## [1] 4e-04

#Different
delta = samp.mean.diff
low = mu0 - delta
high = mu0 + delta
(diff.p = mean(rand.diff>=high) + mean(rand.diff<=low))

## [1] 0
```

- (c) Compute the randomization confidence interval by iterating over values of μ_0 .

```
#####
# Farther
#####

# Lower val
R = 1000
starting.point = mean(dat1$furthest)
lower.val = starting.point
samp.mean = mean(dat1$furthest)
repeat{
  rand = tibble(xbars = rep(NA, R))
  x.shift = dat1$furthest - lower.val
  #Doing the randomization process
  for(i in 1:R){
    curr.samp = x.shift * sample(x = c(-1,1),
                                size = length(x.shift),
                                replace = T)

    rand$xbars[i] = mean(curr.samp)
  }
  #Shifting back
  rand = rand |>
    mutate(xbars = xbars + lower.val)

  #P-value
  delta = abs(samp.mean - lower.val)
  lower = lower.val - delta #Mirror?
  upper = lower.val + delta
  p.val = mean(rand$xbars <= lower) + mean(rand$xbars >= upper)

  if(p.val < .05) {#Should it be .025 or .5?
    break
  }
  else{
    lower.val = lower.val - .0001
  }
}
lower.val

## [1] -0.2534244

#Upper val
upper.val = starting.point
samp.mean = mean(dat1$furthest)
repeat{
  rand = tibble(xbars = rep(NA, 1000))
  x.shift = dat1$furthest - upper.val
  #Doing the randomization process
  for(i in 1:R){
    curr.samp = x.shift * sample(x = c(-1,1),
                                size = length(x.shift),
                                replace = T)

    rand$xbars[i] = mean(curr.samp)
  }
  #Shifting back
  rand = rand |>
    mutate(xbars = xbars + upper.val)

  #P-value
  delta = abs(samp.mean - upper.val)
  lower = upper.val - delta #Mirror?
  upper = upper.val + delta
  p.val = mean(rand$xbars <= lower) + mean(rand$xbars >= upper)
  if(p.val < .05) {#Should it be .025 or .5?
    break
  }
  else{
    upper.val = upper.val + .0001
  }
}
upper.val

## [1] -0.1525244

c(lower.val, upper.val)

## [1] -0.2534244 -0.1525244
```

```
#####
# Closer
#####

# Lower val
R = 1000
starting.point = mean(dat1$closer)
lower.val = starting.point
samp.mean = mean(dat1$closer)
repeat{
  rand = tibble(xbars = rep(NA, R))
  x.shift = dat1$closer - lower.val
  #Doing the randomization process
  for(i in 1:R){
    curr.samp = x.shift * sample(x = c(-1,1),
                                size = length(x.shift),
                                replace = T)

    rand$xbars[i] = mean(curr.samp)
  }
  #Shifting back
  rand = rand |>
    mutate(xbars = xbars + lower.val)

  #P-value
  delta = abs(samp.mean - lower.val)
  lower = lower.val - delta #Mirror?
  upper = lower.val + delta
  p.val = mean(rand$xbars <= lower) + mean(rand$xbars >= upper)

  if(p.val < .05) {#Should it be .025 or .5?
    break
  }
  else{
    lower.val = lower.val - .0001
  }
}
lower.val

## [1] 0.1204231

#Upper val
upper.val = starting.point
samp.mean = mean(dat1$closer)
repeat{
  rand = tibble(xbars = rep(NA, 1000))
  x.shift = dat1$closer - upper.val
  #Doing the randomization process
  for(i in 1:R){
    curr.samp = x.shift * sample(x = c(-1,1),
                                size = length(x.shift),
                                replace = T)

    rand$xbars[i] = mean(curr.samp)
  }
  #Shifting back
  rand = rand |>
    mutate(xbars = xbars + upper.val)

  #P-value
  delta = abs(samp.mean - upper.val)
  lower = upper.val - delta #Mirror?
  upper = upper.val + delta
  p.val = mean(rand$xbars <= lower) + mean(rand$xbars >= upper)
  if(p.val < .05) {#Should it be .025 or .5?
    break
  }
  else{
    upper.val = upper.val + .0001
  }
}
upper.val

## [1] 0.1937231

c(lower.val, upper.val)

## [1] 0.1204231 0.1937231

#####
# Difference
#####
```

```

# Lower val
R = 1000
starting.point = mean(dat1$diff)
lower.val = starting.point
samp.mean = mean(dat1$diff)
repeat{
  rand = tibble(xbars = rep(NA, R))
  x.shift = dat1$diff - lower.val
  #Doing the randomization process
  for(i in 1:R){
    curr.samp = x.shift * sample(x = c(-1,1),
                                size = length(x.shift),
                                replace = T)

    rand$xbars[i] = mean(curr.samp)
  }
  #Shifting back
  rand = rand |>
    mutate(xbars = xbars + lower.val)

  #P-value
  delta = abs(samp.mean - lower.val)
  lower = lower.val - delta #Mirror?
  upper = lower.val + delta
  p.val = mean(rand$xbars <= lower) + mean(rand$xbars >= upper)

  if(p.val < .05) {#Should it be .025 or .5?
    break
  }
  else{
    lower.val = lower.val - .0001
  }
}
lower.val

## [1] 0.2745475

#Upper val
upper.val = starting.point
samp.mean = mean(dat1$diff)
repeat{
  rand = tibble(xbars = rep(NA, 1000))
  x.shift = dat1$diff - upper.val
  #Doing the randomization process
  for(i in 1:R){
    curr.samp = x.shift * sample(x = c(-1,1),
                                size = length(x.shift),
                                replace = T)

    rand$xbars[i] = mean(curr.samp)
  }
  #Shifting back
  rand = rand |>
    mutate(xbars = xbars + upper.val)

  #P-value
  delta = abs(samp.mean - upper.val)
  lower = upper.val - delta #Mirror?
  upper = upper.val + delta
  p.val = mean(rand$xbars <= lower) + mean(rand$xbars >= upper)
  if(p.val < .05) {#Should it be .025 or .5?
    break
  }
  else{
    upper.val = upper.val + .0001
  }
}
upper.val

## [1] 0.4411475

c(lower.val, upper.val)

## [1] 0.2745475 0.4411475

```

Hint: You can “search” for the lower bound from Q_1 and subtracting by 0.0001, and the upper bound using Q_3 and increasing by 0.0001. You will continue until you find the first value for which the two-sided p -value is greater than or equal to 0.05.

4. **Optional Challenge:** In this lab, you performed resampling to approximate the sampling distribution

of the T statistic using

$$T = \frac{\bar{x}_r - 0}{s/\sqrt{n}}.$$

I'm curious whether it is better/worse/similar if we computed the statistics using the sample standard deviation of the resamples (s_r), instead of the original sample (s)

$$T = \frac{\bar{x}_r - 0}{s_r/\sqrt{n}}.$$

- (a) Perform a simulation study to evaluate the Type I error for conducting this hypothesis test both ways.
- (b) Using the same test case(s) as part (a), compute bootstrap confidence intervals and assess their coverage – how often do we ‘capture’ the parameter of interest?

References

Boos, D. D. and Hughes-Oliver, J. M. (2000). How large does n have to be for z and t intervals? *The American Statistician*, 54(2):121–128.